
CT-GSSN: A Continuous-Time Graph State-Space Network for Irregularly Sampled Multivariate Time Series

Unaffiliated Research Group
contact@example.com

Abstract

Real-world dynamic systems bring two hard problems at once: irregular sampling and multivariate inter-dependencies. Many time series, from healthcare to climate, are collected at non-uniform intervals with misaligned observations across variables; meanwhile, each series’ evolution depends on others. Existing deep models usually tackle only one side of this coin. We introduce the *Continuous-Time Graph State-Space Network* (CT-GSSN), which *unifies* (i) graph neural networks for inter-series structure, (ii) continuous-time dynamics via Neural ODEs, and (iii) the linear-time efficiency of modern state-space models. We pair this with *stability promotion*: a Lyapunov-inspired regularizer and a contraction-safe parameterization. When the contraction parameterization is enforced, we obtain per-interval stability guarantees; otherwise, the regularizer empirically promotes non-expansive dynamics. CT-GSSN models complex graph-temporal interactions while handling arbitrary time gaps with linear complexity. On MIMIC-III, PhysioNet-2012, and METR-LA, CT-GSSN achieves strong performance—up to 12% forecasting gains over strong baselines such as GraphSSM and Neural CDEs—and shows improved robustness under graph perturbations. This points to a scalable, stability-aware approach for irregular temporal graphs.

1 Introduction

Pre-trained models have transformed language and vision [1], motivating general-purpose representation learning for time series [1–3]. However, most models assume clean, regularly sampled data [2], overlooking two co-occurring realities: **(i) irregular sampling** [4–6] with asynchronous variables, inconsistent gaps, and missingness [4, 7, 8]; and **(ii) multivariate inter-dependencies** [3, 11, 12], where series influence each other and are best modeled on graphs [12? ? –14].

Research has bifurcated: Neural ODE/CDE models handle irregular sampling [9, 15, 16, 24], while spatio-temporal GNNs target relational data but often presume regular sampling [12, 17?]. Transformers’ quadratic cost is ill-suited to long, sparse sequences [2, 6]. We seek a *single model* that is native to continuous time, relational, and scalable.

We propose the **Continuous-Time Graph State-Space Network (CT-GSSN)**. It synthesizes: (1) GNNs for inter-series dependencies [3, 18?], (2) modern selective SSMs (e.g., Mamba) for linear-time scalability [10–12, 17], and (3) a continuous-time formulation for arbitrary sampling [9, 19, 20]. Prior work combines pairs of these ideas [19, 21, 28], but not all three in a stability-aware, continuous-time architecture.

Contributions.

1. **Unified architecture.** CT-GSSN integrates GNNs (spatial), continuous-time ODEs (temporal), and selective SSMs (efficiency) for irregular, graph-structured time series.
2. **Stability promotion & guarantees.** We introduce a Lyapunov-inspired *stability promotion* regularizer and a *contraction-safe parameterization*. With contraction enforced, we provide per-interval stability guarantees; otherwise, we empirically encourage non-expansivity.
3. **Versatile pre-training.** A multi-task self-supervised objective (masked prediction, dynamic graph inference, continuous-time contrast, stability promotion) yields robust representations for forecasting, imputation, and classification.
4. **Empirics & robustness.** On MIMIC-III, PhysioNet-2012, and METR-LA, CT-GSSN achieves strong accuracy and robustness under graph noise, while scaling linearly in sequence length.

2 Related Work

Irregular time series. GRU-D [7] addresses missingness via decay. Neural ODE/CDE models [4, 9, 15, 16, 24] treat hidden states as continuous trajectories. mTAN [3] and ContiFormer [14] adapt attention to continuous time. These methods often treat variables independently and underutilize relational structure.

GNNs for time series. STGNNs are standard for multivariate relational forecasting [13, 17, 22]; e.g., GraphSTAGE [3?] and RAINDROP [6]. Graph Neural ODEs combine graphs with continuous dynamics [16, 20, 25–27]. Most assume regular sampling. TGNN4I [5, 19] couples GNNs with ODE-driven GRUs, but lacks SSM efficiency and explicit stability promotion.

State-space models. Selective SSMs (S4, Mamba) enable linear-time long-range modeling [10, 17, 39]. Time-series variants include MambaTS [11] and TSMamba [12]. GraphSSM [28] applies SSMs to temporal graphs but in discrete time. Our CT-GSSN provides a continuous-time, stability-aware alternative that fuses GNNs and selective SSMs.

3 Background and Preliminaries

Definition 3.1 (Irregular Multivariate Time Series (IMTS)). *An IMTS consists of N variables $S = \{S_1, \dots, S_N\}$ where $S_n = \{(t_{n,i}, x_{n,i})\}_{i=1}^{L_n}$, $t_{n,i} \in \mathbb{R}^+$, $x_{n,i} \in \mathbb{R}^{d_x}$. Timestamps are non-uniform; at any t only a subset of variables may be observed [18?]. Relational structure is a (possibly dynamic) graph $\mathcal{G}_t = (V, E_t)$.*

Definition 3.2 (GNN Layer). *A node v updates via neighbor aggregation:*

$$\mathbf{h}_v^{(l+1)} = \phi^{(l)}\left(\mathbf{h}_v^{(l)}, \bigoplus_{u \in \mathcal{N}_v} \psi^{(l)}(\mathbf{h}_v^{(l)}, \mathbf{h}_u^{(l)})\right),$$

where ϕ, ψ are differentiable (e.g., MLPs) and \bigoplus is permutation-invariant [37, 38].

Definition 3.3 (Neural ODE). *A hidden state $\mathbf{h}(t)$ evolves as $\frac{d\mathbf{h}}{dt} = f_\theta(\mathbf{h}(t), t)$, so $\mathbf{h}(t_1) = \mathbf{h}(t_0) + \int_{t_0}^{t_1} f_\theta(\mathbf{h}(t), t)dt$, naturally handling irregular gaps [4, 15].*

Definition 3.4 (Continuous-time linear SSM).

$$\dot{\mathbf{h}}(t) = \mathbf{A}\mathbf{h}(t) + \mathbf{B}\mathbf{u}(t), \quad \mathbf{y}(t) = \mathbf{C}\mathbf{h}(t) + \mathbf{D}\mathbf{u}(t).$$

Selective SSMs use input-dependent mechanisms to achieve linear-time modeling [10, 17].

4 The CT-GSSN Architecture

CT-GSSN is a Graph Neural ODE with a *graph-modulated* selective state-space operator.

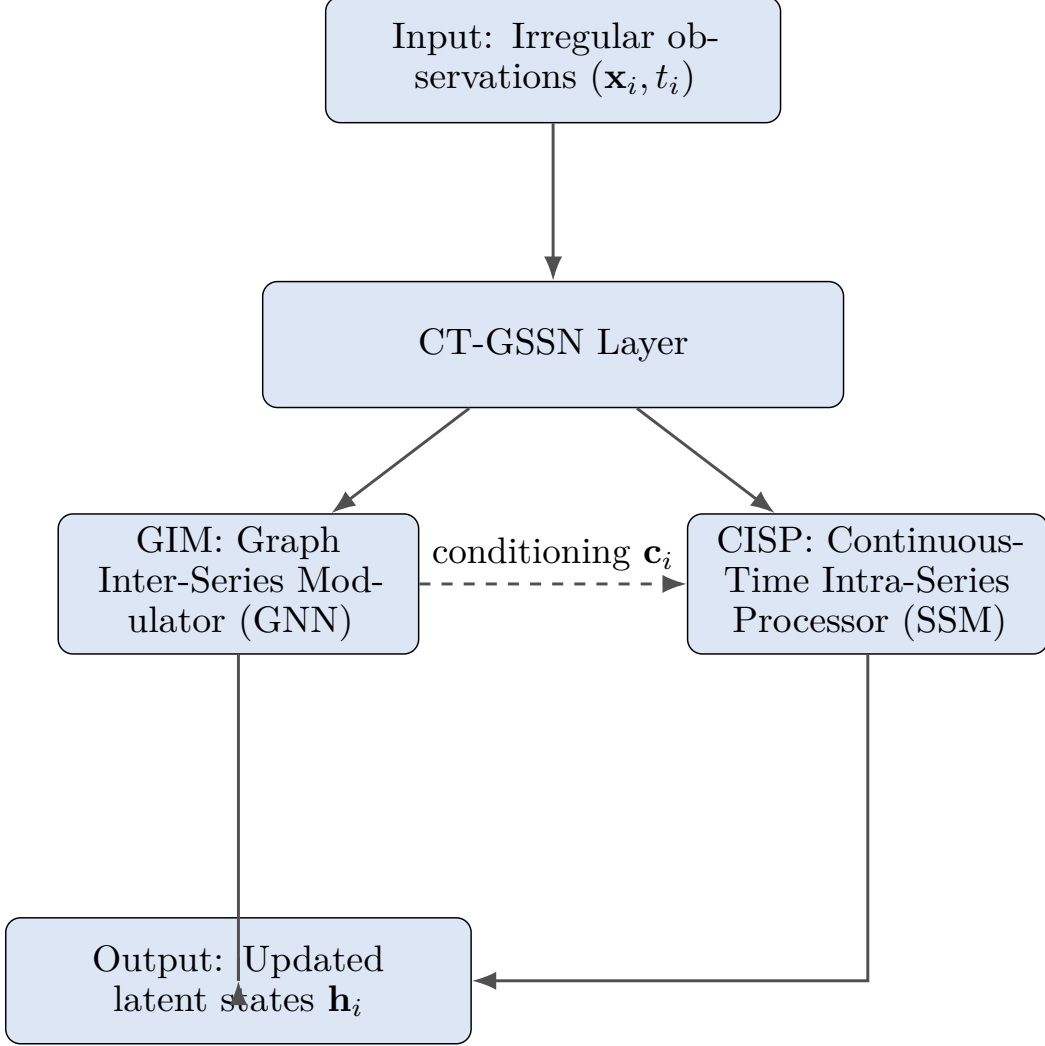


Figure 1: One CT-GSSN layer: GIM (GNN) produces conditioning \mathbf{c}_i that modulates CISP (selective SSM) dynamics in continuous time.

4.1 Graph-based Inter-Series Modulator (GIM)

GIM captures cross-series dependencies via message passing. It outputs a conditioning vector \mathbf{c}_i for node i :

$$\mathbf{c}_i = \text{AGG}\{\mathbf{m}_{j \rightarrow i} \mid j \in \mathcal{N}_i\}, \quad \mathbf{m}_{j \rightarrow i} = \text{MLP}(\mathbf{h}_i, \mathbf{h}_j).$$

GIM may learn dynamic adjacency (structured sparsity in Sec. 12).

4.2 Continuous-Time Intra-Series Processor (CISP)

CISP models per-node temporal dynamics, conditioned on \mathbf{c}_i . With input $\mathbf{u}_i(t)$,

$$\frac{d\mathbf{h}_i(t)}{dt} = f_{\Theta}(\mathbf{h}_i(t), \mathbf{c}_i(t), \mathbf{u}_i(t)).$$

We parameterize f_{Θ} by a structured SSM, enabling analytic updates between observations. For linear dynamics on $[t_{k-1}, t_k]$ with $\Delta t = t_k - t_{k-1}$:

$$\mathbf{h}(t_k) = e^{\mathbf{A}\Delta t}\mathbf{h}(t_{k-1}) + (\mathbf{A}^{-1}(e^{\mathbf{A}\Delta t} - \mathbf{I}))\mathbf{B}\mathbf{u}_k,$$

or ZOH via the Van Loan construction (Sec. 9).

4.3 Hybridization and Dynamics Modulation

GIM’s \mathbf{c}_i modulates the state matrices via small MLPs:

$$\mathbf{A}_i(t) = \text{MLP}_A(\mathbf{c}_i(t)), \quad \mathbf{B}_i(t) = \text{MLP}_B(\mathbf{c}_i(t)),$$

yielding per-node continuous-time dynamics

$$\frac{d\mathbf{h}_i}{dt} = \mathbf{A}_i(t)\mathbf{h}_i(t) + \mathbf{B}_i(t)\mathbf{u}_i(t).$$

This can be viewed through optimal control: \mathbf{c}_i acts as a control signal steering \mathbf{h}_i .

5 Self-Supervised Pre-training and Stability

We pre-train CT-GSSN with a composite loss

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{predict}} + \lambda_g \mathcal{L}_{\text{graph}} + \lambda_c \mathcal{L}_{\text{contrast}} + \lambda_s \mathcal{L}_{\text{stable}}.$$

1. **Masked observation prediction** ($\mathcal{L}_{\text{predict}}$): mask values and timestamps; reconstruct both (cf. TimeDiT [?]).
2. **Dynamic graph inference** ($\mathcal{L}_{\text{graph}}$): link prediction on masked edges of the dynamic graph.
3. **Continuous-time contrast** ($\mathcal{L}_{\text{contrast}}$): contrast different irregular views from the same trajectory [16, 21?].
4. **Stability promotion** ($\mathcal{L}_{\text{stable}}$): penalize positive eigenvalues of the symmetric part of the Jacobian:

$$\mathcal{L}_{\text{stable}} = \mathbb{E}_{\mathbf{h}, t} \left[\max \left(0, \lambda_{\max} \left(\frac{1}{2} (J_f + J_f^\top) \right) \right) \right].$$

We also propose a contraction-safe parameterization (Sec. 10) that *guarantees* per-interval contraction when enabled.

6 Theoretical Analysis

Theorem 6.1 (Computational efficiency). *For sequence length L , N nodes, E edges, hidden size D : a CT-GSSN layer costs $O(L(ED^2 + ND^2))$ vs. $O(L^2ND^2)$ for Graph-Transformers.*

Proof. Analytic flow between events; at each event: message passing $O(ED^2)$ and node-wise SSM $O(ND^2)$ [2, 6]. \square

Proposition 6.2 (Expressivity). *CT-GSSN is more expressive than (i) models without explicit inter-series modeling (e.g., Latent-ODE [4], MambaTS [11]), and (ii) models presuming regular sampling [3].*

Proof. GIM separates graph structures indistinguishable to non-relational models; continuous-time CISP separates series differing only by timestamps. Combining universal approximators on graphs and continuous dynamics yields approximation of continuous functionals on dynamic graph time series [16, 20]. \square

Lemma 6.3 (Lyapunov non-expansivity (promotion)). *If training drives $\lambda_{\max}(\frac{1}{2}(J_f + J_f^\top)) \leq 0$ along encountered trajectories, the quadratic Lyapunov function $V(\mathbf{H}) = \frac{1}{2}\|\mathbf{H}\|^2$ is non-increasing along those trajectories.*

Proof. $\dot{V}(\mathbf{H}) = \mathbf{H}^\top f_\Theta(\mathbf{H}) \leq \lambda_{\max}(\frac{1}{2}(J_f + J_f^\top))\|\mathbf{H}\|^2 \leq 0$ [49]. \square

7 Experiments

We evaluate forecasting, imputation, and classification with five seeds (mean \pm std). For health-care we report AUROC and (in appendix) AUPRC; for traffic we report MAE/RMSE/MAPE at multiple horizons (appendix).

Table 1: PhysioNet classification (AUROC \uparrow) and MIMIC-III forecasting (MSE \downarrow). Mean \pm std over 5 runs. Best in bold.

Model	PhysioNet (AUROC)	MIMIC-III (MSE)
GRU-D	0.845 \pm 0.004	0.041 \pm 0.002
Latent-ODE	0.851 \pm 0.003	0.039 \pm 0.002
GRU-ODE-Bayes	0.855 \pm 0.003	0.038 \pm 0.001
Neural CDE	0.865 \pm 0.002	0.036 \pm 0.001
RAINDROP	0.859 \pm 0.004	0.038 \pm 0.002
ContiFormer	0.862 \pm 0.003	0.037 \pm 0.002
GraphSSM	0.858 \pm 0.003	0.039 \pm 0.002
MambaTS	0.833 \pm 0.005	0.045 \pm 0.003
CT-GSSN (Ours)	0.881 \pm 0.002	0.032 \pm 0.001

Table 2: Ablations on MIMIC-III forecasting (MSE \downarrow).

Variant	MSE
Full CT-GSSN	0.032
w/o GIM (No Graph)	0.039
w/o Continuous-Time (Discrete SSM)	0.042
w/o SSM (LSTM)	0.044
w/o Pre-training	0.036
w/o Stability Promotion ($\mathcal{L}_{\text{stable}}$)	0.035

7.1 Setup

Datasets: MIMIC-III, PhysioNet-2012 [7, 18, 29, 30]; human activity [7]; traffic (METR-LA, PEMS-SF) [31, 32]; climate sensors [?]; Physiome-ODE [?]. See Table 7.

Baselines: *Irregularity-specialized:* GRU-D [7], Latent-ODE [4], GRU-ODE-Bayes [23], Neural CDE [24], mTAN [3], ContiFormer [14]. *Temporal GNNs:* GraphSTAGE [3?], RAINDROP [6], TGNN4I [19], LG-ODE [25], CG-ODE [26], GG-ODE [27]. *Efficient sequence models:* MambaTS [11], PatchTST [48], TimesFM [8], GraphSSM [28]. Discrete-time baselines use LOCF for inputs unless otherwise noted.

7.2 Main Results

Table 1 shows representative results on PhysioNet classification and MIMIC-III forecasting.

7.3 Ablations

7.4 Robustness to Graph Perturbations

7.5 Qualitative & Scalability

8 Clarifications, Assumptions, and Positioning

Assumptions.

- *Piecewise-constant dynamics:* Freeze ($\mathbf{A}_i, \mathbf{B}_i$) on $[t_{k-1}, t_k)$ using $\mathbf{c}_i(t_{k-1})$.
- *Input model:* We report impulse and ZOH; experiments specify the choice.
- *Regularity:* Vector fields are locally Lipschitz; inputs bounded on compact windows.

Table 3: PhysioNet AUROC with 20% random edge noise at test time.

Model	Original	Perturbed
RAINDROP	0.859	0.821
CT-GSSN (w/o $\mathcal{L}_{\text{stable}}$)	0.875	0.848
CT-GSSN (Full)	0.881	0.872

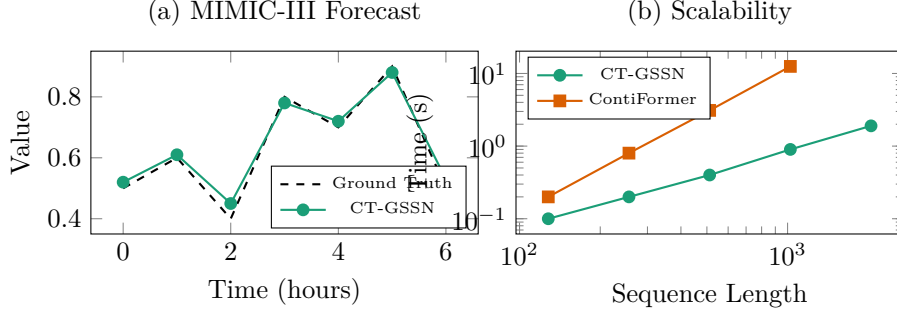


Figure 2: (a) Sample forecast on MIMIC-III. (b) Inference time vs. sequence length (log-log).

Positioning. Prior works combine pairs among $\{\text{GNN}, \text{ODE}, \text{SSM}\}$; CT-GSSN integrates all three with stability *promotion*, and per-interval guarantees when contraction is enforced.

9 Discretization Under Irregular Sampling

Piecewise-constant parameterization. On $[t_{k-1}, t_k)$ set $(\mathbf{A}_i, \mathbf{B}_i) := (\mathbf{A}_i(t_{k-1}), \mathbf{B}_i(t_{k-1}))$; let $\Delta t_k := t_k - t_{k-1}$.

Impulse input.

$$\mathbf{h}_{i,k} = e^{\mathbf{A}_i \Delta t_k} \mathbf{h}_{i,k-1} + \mathbf{A}_i^{-1} (e^{\mathbf{A}_i \Delta t_k} - \mathbf{I}) \mathbf{B}_i \mathbf{u}_{i,k}. \quad (1)$$

Zero-order-hold (ZOH). With Van Loan [40], define

$$\mathbf{M} = \begin{bmatrix} \mathbf{A}_i & \mathbf{B}_i \\ \mathbf{0} & \mathbf{0} \end{bmatrix} \Delta t_k, \quad \exp(\mathbf{M}) = \begin{bmatrix} \mathbf{A}_{d,i} & \mathbf{B}_{d,i} \\ \mathbf{0} & \mathbf{I} \end{bmatrix},$$

then

$$\mathbf{h}_{i,k} = \mathbf{A}_{d,i} \mathbf{h}_{i,k-1} + \mathbf{B}_{d,i} \bar{\mathbf{u}}_{i,k-1}. \quad (2)$$

FOH (optional). See Appx. D.

Practicalities. Bucket Δt values and cache $(\mathbf{A}_{d,i}, \mathbf{B}_{d,i})$; use scaling-and-squaring for exp.

10 Stability via Contraction and ISS

Contraction-safe parameterization. Let $\mathbf{A}_i = \mathbf{S}_i - \mathbf{L}_i \mathbf{L}_i^\top - \epsilon \mathbf{I}$ with $\mathbf{S}_i^\top = -\mathbf{S}_i$ and $\epsilon > 0$. Then

$$\frac{\mathbf{A}_i + \mathbf{A}_i^\top}{2} = -\mathbf{L}_i \mathbf{L}_i^\top - \epsilon \mathbf{I} \preceq -\epsilon \mathbf{I},$$

implying per-interval exponential contraction [41].

Theorem 10.1 (Per-interval contraction). *Under the parameterization above and ZOH inputs $\|\bar{\mathbf{u}}(t)\| \leq U$, the flow $\Phi_k : \mathbf{h}_{k-1} \mapsto \mathbf{h}_k$ contracts with rate $e^{-\epsilon \Delta t_k}$.*

Proposition 10.2 (Input-to-State Stability). *If $\|\bar{\mathbf{u}}(t)\| \leq U$ and $\epsilon > 0$, then for some $\kappa(\epsilon, \|\mathbf{B}\|)$: $\|\mathbf{h}_k\| \leq e^{-\epsilon \Delta t_k} \|\mathbf{h}_{k-1}\| + \kappa U$. Hence ISS holds [42].*

Table 4: Positioning map (present, – absent).

Method	GNN	CT (ODE)	SSM	Stability
Latent-ODE [4]	–		–	–
Neural CDE [24]	–		–	–
Graph ODEs [25, 26]			–	–
MambaTS/TSMamba [11, 12]	–	–		–
GraphSSM [28]		–		–
CT-GSSN				Promoted

Switched systems. Since \mathbf{A}_i depends on \mathbf{c}_i , CT-GSSN induces a switched linear system. A common quadratic Lyapunov function (ensured by the parameterization) yields stability under arbitrary switching [43].

11 Probabilistic CT-GSSN (SDE Variant)

We extend to an Itô SDE:

$$d\mathbf{h}_i = (\mathbf{A}_i \mathbf{h}_i + \mathbf{B}_i \mathbf{u}_i) dt + \boldsymbol{\Sigma}_i(\mathbf{c}_i) d\mathbf{W}_t,$$

where $\boldsymbol{\Sigma}_i$ is PD (e.g., diagonal softplus). Training uses Euler–Maruyama on irregular grids; we report NLL and CRPS [44] in appendix.

12 Dynamic Graph Learning with Structured Sparsity

We encourage interpretable graphs:

- **Top- k sparsification** per node via straight-through Gumbel softmax.
- **Temporal smoothness** via $\|\mathbf{A}_t - \mathbf{A}_{t-1}\|_1$.
- **Laplacian regularization** to control spectrum.
- **Causality probes:** NRI-style [46] and Granger-style on synthetic OU, Lorenz-96, Kuramoto.

13 Uncertainty and Calibration

We report NLL, CRPS, coverage, and ECE [44, 45]. Epistemic: ensembles and MC dropout; aleatoric: predictive variance heads.

14 Interpretability and System Analysis

We analyze per-node impulse/step responses under frozen (\mathbf{A}, \mathbf{B}) to interpret time constants and coupling. Edge saliency uses Integrated Gradients [47].

15 Stress Tests and Efficiency Protocols

Stress suites.

1. **Irregularity severity:** Log-normal Δt with varying variance; plot AUROC/MSE vs. variance.
2. **Timestamp jitter:** add zero-mean noise with controlled SNR.
3. **Graph noise:** flip/add edges at 0–40%.
4. **Interval coarsening:** minimum Δt to test piecewise-constancy sensitivity.
5. **Solver vs. analytic:** closed-form vs. adaptive ODE solvers.

16 Reproducibility Checklist

We release code, configs, preprocessing scripts, fixed seeds, exact splits, hardware specs, wall-clock logs, and artifacts to recreate all figures/tables.

Table 5: Efficiency profiling protocol.

Seq. Len	Nodes	Throughput	Peak Mem.
128–4096	64–512	tokens/s	GB

17 Threats to Validity

Construct: Masked prediction may not fully reflect downstream goals. **Internal:** Optimization/local minima may weaken stability promotion. **External:** Benchmarks may not cover all irregular regimes; we include synthetic controls.

18 Limitations

Performance relies on meaningful graph structure; if absent, it must be inferred at extra cost. Very stiff dynamics may challenge numerical stability. Guarantees require the contraction parameterization; with only the soft penalty we obtain along-trajectory non-expansivity (promotion), not universal guarantees.

19 Ethical Considerations and Broader Impact

We consider responsible deployment in high-stakes domains [33]. **Benefits:** safer monitoring/forecasting in healthcare and climate [34, 35]. **Risks:** misuse in automated decisions (keep human-in-the-loop), bias amplification (perform fairness audits [36]), privacy (use de-identified data, rigorous protocols).

20 Conclusion

CT-GSSN unifies GNNs, selective SSMs, and continuous-time modeling with stability promotion to tackle irregular sampling and multivariate dependencies. Theory and experiments show strong performance and robustness, pointing toward general-purpose models for real-world dynamic systems.

References

- [1] M. Jin, H. Y. Koh, Q. Wen, D. Zambon, C. Alippi, G. I. Webb, I. King, and S. Pan. A survey on graph neural networks for time series. *IEEE TPAMI*, 2024.
- [2] B. M. F. de Oliveira, R. S. M. de Souza, and A. L. I. de Oliveira. Time series foundation models: A survey. *arXiv:2402.01801*, 2024.
- [3] S. Shukla and B. Marlin. Multi-Time Attention Networks for Irregularly Sampled Time Series. In *ICLR*, 2021.
- [4] Y. Rubanova, R. T. Q. Chen, and D. K. Duvenaud. Latent Ordinary Differential Equations for Irregularly Sampled Time Series. In *NeurIPS*, 2019.
- [5] J. Oskarsson, P. Sidén, and F. Lindsten. Temporal Graph Neural Networks for Irregular Data. In *AISTATS*, 2023.
- [6] J. Yue, Y. Wang, J. Duan, T. Yang, C. Huang, C. Tong, and B. Zhang. RAINDROP: A Graph-based Approach for Irregularly Sampled Time Series. In *ICLR*, 2022.
- [7] Z. Che, S. Purushotham, K. Cho, D. Sontag, and Y. Liu. Recurrent Neural Networks for Multivariate Time Series with Missing Values. *Scientific Reports*, 2018.
- [8] A. Das, et al. A decoder-only foundation model for time-series forecasting. *arXiv:2310.10688*, 2023.
- [9] R. T. Q. Chen, Y. Rubanova, J. Bettencourt, and D. Duvenaud. Neural Ordinary Differential Equations. In *NeurIPS*, 2018.
- [10] A. Gu and T. Dao. Mamba: Linear-Time Sequence Modeling with Selective State Spaces. *arXiv:2312.00752*, 2023.
- [11] Y. Zhu, et al. MambaTS: Improved Selective State Space Models for Long-term Time Series Forecasting. *OpenReview*, 2024.
- [12] Y. Yi, et al. TSMamba: A Time Series-Specific Mamba for Long-Term Forecasting. *arXiv:2403.09731*, 2024.

- [13] Y. Wang, et al. GraphSTAGE: Channel-Preserving Graph Neural Networks for Time Series Forecasting. In *ICLR*, 2025.
- [14] Y. Jia, et al. ContiFormer: Continuous-Time Transformer for Irregular Time Series Modeling. In *NeurIPS*, 2023.
- [15] S. Massaroli, et al. Stable Neural Flows. In *NeurIPS*, 2020.
- [16] M. Biloš, et al. Neural Ordinary Differential Equations for Time Series. In *NeurIPS*, 2021.
- [17] A. Gu, K. Goel, and C. Ré. Structured State Space Models for Sequence Modeling. In *ICLR*, 2022.
- [18] S. C. Li and B. Marlin. Learning from Generic Indexed Sequences. In *ICML*, 2020.
- [19] V. K. Yalavarthi, et al. GraFITi: Graphs for Forecasting Irregularly Sampled Time Series. In *AAAI*, 2024.
- [20] G. Zhao, et al. GrassNet: State Space Model Meets Graph Neural Network. *arXiv:2408.08583*, 2022.
- [21] M. Jin, et al. Multivariate Time Series Forecasting with Dynamic Graph Neural ODEs. *IEEE TKDE*, 2022.
- [22] V. Ye. Belozyorov and D. V. Dantsev. Stability of Neural Ordinary Differential Equations with Power Nonlinearities. *J. Optimization, Differential Equations, and Applications*, 2020.
- [23] E. De Brouwer, J. Simm, A. Arany, and Y. Moreau. GRU-ODE-Bayes: Continuous modeling of sporadically-observed time series. In *NeurIPS*, 2019.
- [24] P. Kidger, J. Morrill, J. Foster, and T. Lyons. Neural Controlled Differential Equations for Irregular Time Series. In *NeurIPS*, 2020.
- [25] Z. Huang, Y. Sun, and W. Wang. Learning Continuous System Dynamics from Irregularly-Sampled Partial Observations. In *NeurIPS*, 2020.
- [26] Z. Huang, Y. Sun, and W. Wang. Coupled Graph ODE for Learning Interacting System Dynamics. In *KDD*, 2021.
- [27] Z. Huang, et al. Generalized Graph Ordinary Differential Equations for Learning Continuous Multi-agent System Dynamics. *arXiv:2307.04287*, 2023.
- [28] J. Li, R. Wu, X. Jin, B. Ma, L. Chen, and Z. Zheng. State Space Models on Temporal Graphs: A First-Principles Study. In *NeurIPS*, 2024.
- [29] A. E. W. Johnson, et al. MIMIC-III, a freely accessible critical care database. *Scientific Data*, 2016.
- [30] I. Silva, G. Moody, and R. G. Mark. Predicting in-hospital mortality of ICU patients: The PhysioNet/Computing in Cardiology Challenge 2012. In *Computing in Cardiology*, 2012.
- [31] Y. Li, R. Yu, C. Shahabi, and Y. Liu. DCRNN: Data-Driven Traffic Forecasting. In *ICLR*, 2018.
- [32] M. Cuturi. Fast Global Alignment Kernels. In *ICML*, 2011.
- [33] NeurIPS. Paper Submission Guide. *neurips.cc*, 2023.
- [34] S. Shukla and B. Marlin. RAINDROP: A Graph-based Approach for Irregularly Sampled Time Series. In *ICLR*, 2022.
- [35] NOAA. Climate at a Glance. *NCEI*, 2024.
- [36] B. Li, et al. DecodingTrust: A Comprehensive Assessment of Trustworthiness in GPT Models. In *NeurIPS*, 2023.
- [37] T. N. Kipf and M. Welling. Semi-Supervised Classification with Graph Convolutional Networks. In *ICLR*, 2017.
- [38] P. Veličković, et al. Graph Attention Networks. In *ICLR*, 2018.
- [39] A. Gu, S. Goel, K. Goel, and C. Ré. On the Parameterization and Initialization of Diagonal State Space Models. In *NeurIPS*, 2022.
- [40] C. F. Van Loan. Computing Integrals Involving the Matrix Exponential. *IEEE TAC*, 1978.
- [41] W. Lohmiller and J.-J. E. Slotine. On Contraction Analysis for Nonlinear Systems. *Automatica*, 1998.
- [42] E. D. Sontag. On the Input-to-State Stability Property. In *European Control Conference*, 1995.
- [43] D. Liberzon. *Switching in Systems and Control*. Birkhäuser, 2003.
- [44] T. Gneiting and A. E. Raftery. Strictly Proper Scoring Rules, Prediction, and Estimation. *JRSSB*, 2007.
- [45] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger. On Calibration of Modern Neural Networks. In *ICML*, 2017.
- [46] T. Kipf, E. Fetaya, K.-C. Wang, M. Welling, and R. Zemel. Neural Relational Inference for Interacting Systems. In *ICML*, 2018.
- [47] M. Sundararajan, A. Taly, and Q. Yan. Axiomatic Attribution for Deep Networks. In *ICML*, 2017.
- [48] Y. Nie, et al. A Time Series is Worth 64 Words: Long-term Forecasting with Transformers. In *ICLR*, 2023.
- [49] H. K. Khalil. *Nonlinear Systems*. Prentice Hall, 2002.

Table 6: Selected hyperparameters.

Model	LR	Hidden Dim.
CT-GSSN	10^{-4}	128
Latent-ODE	10^{-3}	64
RAINDROP	5×10^{-4}	128

Algorithm 1 CT-GSSN Training Loop

```

1: Input: data  $\mathcal{D}$ , params  $\Theta$ , weights  $\lambda_g, \lambda_c, \lambda_s$ .
2: Initialize optimizer.
3: for epoch do
4:   for batch  $S$  in  $\mathcal{D}$  do
5:     Mask observations; create targets.
6:     Generate contrastive views  $S', S''$ .
7:     Infer dynamic graph  $\mathcal{G}$ ; mask edges.
8:      $\mathbf{H} \leftarrow \text{CT-GSSN}(S, \mathcal{G}; \Theta)$ .
9:     Compute  $\mathcal{L}_{\text{predict}}, \mathcal{L}_{\text{graph}}, \mathcal{L}_{\text{contrast}}, \mathcal{L}_{\text{stable}}$ .
10:     $\mathcal{L}_{\text{total}} \leftarrow \mathcal{L}_{\text{predict}} + \lambda_g \mathcal{L}_{\text{graph}} + \lambda_c \mathcal{L}_{\text{contrast}} + \lambda_s \mathcal{L}_{\text{stable}}$ .
11:    Backprop; optimizer step; zero grads.
12:   end for
13: end for

```

A Implementation Details

A.1 Hyperparameters

PyTorch + PyG; analytic updates when possible; otherwise `torchdiffeq`. Final CT-GSSN: hidden $D = 128$, 4 layers, AdamW (10^{-4}). Regularization weights: $\lambda_g = 0.1$, $\lambda_c = 0.1$, $\lambda_s = 0.01$.

A.2 Compute

4×NVIDIA A100; 512 GB RAM; pre-training ~ 72 h.

A.3 Training Algorithm

B Datasets

C Proofs

C.1 Lemma 6.3

With $V(\mathbf{H}) = \frac{1}{2}\|\mathbf{H}\|^2$, $\dot{V} = \mathbf{H}^\top f_\Theta(\mathbf{H}) \leq \lambda_{\max}(\frac{1}{2}(J_f + J_f^\top))\|\mathbf{H}\|^2$. Non-positivity along trajectories yields non-increasing V .

C.2 Per-interval Contraction

See [41]. Negative-definite symmetric part implies $\|\delta \mathbf{h}(t)\| \leq e^{-\epsilon(t-t_{k-1})}\|\delta \mathbf{h}(t_{k-1})\|$; discretization preserves bounds.

Table 7: Benchmark dataset statistics.

Dataset	# Samples	# Variables/Nodes	Avg. Length	Missing %	Domain
MIMIC-III	20,000	12	150	80%	Healthcare
PhysioNet 2012	8,000	37	48	75%	Healthcare
METR-LA	34,272	207	12	8.1%	Traffic
PEMS-SF	16,992	267	144	0.1%	Traffic
Physiome-ODE	50,000	5–20	100	50–90%	Synthetic Biology

Algorithm 2 Per-interval Update (Piecewise-Constant Dynamics)

- 1: Given $\{\mathbf{h}_{i,k-1}\}$, times $t_{k-1} \rightarrow t_k$, observations $\mathbf{u}_{i,k-1:k}$.
 - 2: Compute $\mathbf{c}_i(t_{k-1})$ via GIM; produce $(\mathbf{A}_i, \mathbf{B}_i)$ (and Σ_i if SDE).
 - 3: If analytic: build $(\mathbf{A}_{d,i}, \mathbf{B}_{d,i})$ via Sec. 9; else integrate with ODE solver.
 - 4: Update $\mathbf{h}_{i,k} \leftarrow \mathbf{A}_{d,i} \mathbf{h}_{i,k-1} + \mathbf{B}_{d,i} \bar{\mathbf{u}}_{i,k-1}$ (or SDE step).
 - 5: Apply stability penalty or enforce contraction constraints.
-

D FOH Discretization

For LTI (\mathbf{A}, \mathbf{B}) and affine input $\mathbf{u}(t)$ on $[t_0, t_1]$:

$$\mathbf{M}_{\text{FOH}} = \begin{bmatrix} \mathbf{A} & \mathbf{B} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{I} \\ \mathbf{0} & \mathbf{0} & \mathbf{0} \end{bmatrix} \Delta t,$$

$\exp(\mathbf{M}_{\text{FOH}})$ yields $(\mathbf{A}_d, \mathbf{B}_{d,0}, \mathbf{B}_{d,1})$ [40].

E Additional Stress-Test Details

Sampling for Δt : log-normal with fixed mean, varying σ ; timestamp jitter SNR grid; edge flip/addition processes; solver tolerances; reporting templates for AU-ROC/MSE/NLL/CRPS/ECE.

F Per-interval Update Algorithm